

Task 3 Paper Review and Reproducibility Note

Fast High-Resolution Image Synthesis with Latent Adversarial Diffusion Distillation
Axel Sauer, Frederic Boesel, Tim Dockhorn, Andreas Blattmann, Patrick Esser, Robin Rombach
arXiv:2403.12015 (2024)

1. Problem and Motivation

Diffusion models produce high-quality images, but they are often slow at inference time because generation usually requires many sequential denoising steps. This creates a practical bottleneck for interactive image synthesis, editing, and deployment. The paper targets this speed-quality trade-off directly. Its core question is whether a strong multi-step text-to-image diffusion model can be distilled into a much faster generator without losing too much quality. This is a practically important problem because a method that preserves useful visual fidelity at very low step counts could make diffusion systems more usable in real-world settings.

2. Core Method

The paper proposes Latent Adversarial Diffusion Distillation (LADD). The key idea is to distill a stronger diffusion teacher into a fast student model while moving the adversarial supervision into a latent feature space rather than relying only on pixel-level matching. Compared with earlier adversarial diffusion distillation methods, LADD uses generative features from pretrained latent diffusion models. This design is intended to simplify optimization and improve visual quality, especially for high-resolution and multi-aspect-ratio synthesis. In broad terms, the method keeps the low-step generation goal of earlier distillation approaches but uses a more suitable representation for training supervision.

3. Data, Evaluation, and Main Findings

The paper studies high-resolution text-to-image generation and reports results on a distilled system derived from Stable Diffusion 3. From an evaluation perspective, the paper focuses on standard text-to-image generation quality and alignment assessment, together with qualitative visual comparisons under very small step counts. This is central to the paper because the goal is not only to reduce inference cost, but also to preserve useful image quality after aggressive distillation. In addition to qualitative comparisons, the paper reports benchmark-based evidence to support the claim that the distilled model remains competitive under aggressive step reduction. In that sense, the evaluation is centered on standard text-to-image generation quality and alignment assessments used in the paper, together with direct visual comparisons across outputs. The main reported outcome is that LADD enables SD3-Turbo, a model that can match strong contemporary text-to-image performance using only four unguided sampling steps. The paper also discusses scaling behavior and shows applications such as image editing and inpainting. The main contribution is therefore not only a new training objective, but also evidence that fast distilled diffusion can remain visually competitive under severe step reduction.

4. Two Strengths

First, the paper addresses a highly relevant systems-level limitation of diffusion models: inference cost. This makes the contribution practical rather than purely theoretical. Second, the method is conceptually well motivated. By using latent generative features instead of fixed pixel-space adversarial supervision, the paper proposes a training signal that is more naturally aligned with latent diffusion pipelines and high-resolution image synthesis.

5. Two Limitations and One Reproducibility Challenge

One limitation is that the full training setup is expensive and difficult to reproduce in a course environment. Recreating the method would likely require large models, significant compute, and access to the exact checkpoints or training configuration. A second limitation is that the paper's strongest claims are supported by large-scale qualitative and benchmark evidence, which makes a small independent reproduction necessarily partial. The main reproducibility challenge is therefore not understanding the idea, but

reproducing the original training and evaluation pipeline faithfully enough to verify the full reported gains. More specifically, the main barriers are the scale of the original training setup, the dependence on large pretrained diffusion components, and the difficulty of matching the original compute budget and evaluation pipeline in a course setting. For that reason, the reproducibility component here is framed as a structured proxy experiment rather than a claim of full paper-level reproduction.

6. Small Reproducibility Component

Because full paper-scale reproduction is not feasible here, I include a compact proxy experiment that tests one central claim behind the paper: whether very small inference-step counts can still produce coherent images, and how image quality changes as the number of steps increases. The figure below compares outputs generated with 1, 2, and 4 steps using the same city-at-sunset prompt.




Figure 1. Proxy step-count comparison from the notebook: 1-step, 2-step, and 4-step outputs for the same prompt.

These outputs are included because a table of settings alone would not show whether the generated scene remains structurally recognizable under aggressive step compression. The 1-step image already captures the global concept of a lit city skyline at sunset, which is important because it shows that extreme speed-up can still preserve coarse semantic structure. At 2 steps, the buildings, windows, and lighting patterns become more regular, and the geometry of the skyline is visibly more stable. At 4 steps, the image is the most refined among the three: building edges are cleaner, window patterns are more consistent, and the overall output appears more polished. This supports the paper's broader motivation that fast distilled diffusion can produce usable images at very low step counts, while also making the remaining quality trade-off visible.

7. Follow-up Experiment

A concrete next step would be to extend the qualitative step-count comparison with a small quantitative study. For example, I would compare runtime together with an image-quality proxy across 1, 2, 4, and possibly 8 steps, and then analyze where the diminishing-return region begins. That would strengthen the review because it would connect the visual evidence more directly to the speed-quality trade-off that motivates the paper.

References

Sauer, A., Boesel, F., Dockhorn, T., Blattmann, A., Esser, P., & Rombach, R. (2024). *Fast High-Resolution Image Synthesis with Latent Adversarial Diffusion Distillation*. arXiv:2403.12015.

Hugging Face Diffusers Documentation. Consulted for practical pipeline usage and inference-step configuration in the notebook component.

Model and pipeline resources related to Stable Diffusion and fast distilled diffusion systems. Used as supporting implementation references for the compact proxy experiment.